

Deep Learning

Lecture 05: Recurrent Neural Networks and Transformers

Wanxiang Che

Recurrent Neural Networks

Language Models


- A language model computes a probability for a sequence of word:
 $P(w_1, \dots w_n)$ or predicts a probability for the next word:
 $P(w_{n+1} | w_1, \dots w_n)$
- Useful for machine translation, speech recognition, and so on
 - Word ordering
 - $P(\text{the cat is small}) > P(\text{small the is cat})$
 - Word choice
 - $P(\text{there are } \mathbf{four} \text{ cats}) > P(\text{there are } \mathbf{for} \text{ cats})$

Traditional Language Models

- An incorrect but necessary **Markov** assumption!
 - Probability is usually conditioned on window of $n-1$ previous words
 - $P(w_1, \dots, w_m) = \prod_{i=1}^m P(w_i | w_1, \dots, w_{i-1}) \approx \prod_{i=1}^m P(w_i | w_{i-(n-1)}, \dots, w_{i-1})$
- How to estimate probabilities
 - $p(w_2 | w_1) = \frac{\text{count}(w_1, w_2)}{\text{count}(w_1)}$ $p(w_3 | w_1, w_2) = \frac{\text{count}(w_1, w_2, w_3)}{\text{count}(w_1, w_2)}$
- Performance improves with keeping around higher n-grams counts and doing **smoothing**, such as backoff (e.g. if 4-gram not found, try 3-gram, etc)
- Disadvantages
 - There are A LOT of n-grams!
 - Cannot see too long history
 - $P(\text{坐/作 了一整天的 火车/作业})$


Recurrent Neural Networks (RNNs)

- Condition the neural network on all previous inputs
- RAM requirement only scales with number of inputs



Recurrent Neural Networks (RNNs)

- At a single time step t
 - $h_t = \tanh(W^1 h_{t-1} + W^2 x_t)$
 - $\hat{y}_t = \text{softmax}(W^3 h_t)$



Recurrent Neural Network Language Model (RNNLM)


- Modeling the infinite context (Mikolov et al. 2010-2013)

- Compute:

$$\mathbf{h}(t) = f(\mathbf{U}\mathbf{w}(t) + \mathbf{W}\mathbf{h}(t-1))$$


$$\mathbf{y}(t) = g(\mathbf{V}\mathbf{h}(t))$$

- \mathbf{U} is the Embedding Matrix



BackPropagation Through Time (BPTT)

- Total error is the sum of each error at time step t
 - $\frac{\partial E}{\partial W} = \sum_{t=1}^T \frac{\partial E_t}{\partial W}$
- $\frac{\partial E_t}{\partial W^3} = \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial W^3}$ is easy to be calculated
- But to calculate $\frac{\partial E_t}{\partial W^1} = \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial W^1}$ is hard (also for W^2)
- Because $h_t = \tanh(W^1 h_{t-1} + W^2 x_t)$ depends on h_{t-1} , which depends on W^1 and h_{t-2} , and so on.



BackPropagation Through Time (BPTT)

- Use the same as the backpropagation algorithm as we use in deep feedforward NN, but summing up the gradients for W^1
- BPTT is just a fancy name for standard backpropagation on an unrolled RNN

- $$\frac{\partial E}{\partial W^1} = \sum_{t=1}^T \frac{\partial E_t}{\partial W^1}$$
- $$\frac{\partial E_t}{\partial W^1} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W^1}$$

The vanishing gradient problem

- $\frac{\partial E_t}{\partial W^1} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W^1}$, $h_t = \tanh(W^1 h_{t-1} + W^2 x_t)$
- $\frac{\partial h_t}{\partial h_k} = \prod_{j=k+1}^t \frac{\partial h_j}{\partial h_{j-1}} = \prod_{j=k+1}^t W^1 \text{diag}[\tanh'(\dots)]$
- $\left\| \frac{\partial h_t}{\partial h_{t-1}} \right\| \leq \gamma \|W^1\| \leq \gamma \lambda_1$
 - where γ is bound $\|\text{diag}[\tanh'(\dots)]\|$, λ_1 is the largest singular value of W^1
- $\left\| \frac{\partial h_t}{\partial h_k} \right\| \leq (\gamma \lambda_1)^{t-k} \rightarrow 0$, if $\gamma \lambda_1 < 1$
- This can become very small or very large quickly \rightarrow Vanishing or exploding gradient
 - Trick for exploding gradient: clipping trick (set a threshold)

A “solution”

- Intuition
 - Ensure $\gamma\lambda_1 \geq 1 \rightarrow$ to prevent vanishing gradients
- So ...
 - Proper initialization of the W
 - To use ReLU instead of tanh or sigmoid activation functions

A better “solution”

- Recall the original transition equation

- $h_t = \tanh(W^1 h_{t-1} + W^2 x_t)$

- We can instead update the state **additively**


- $u_t = \tanh(W^1 h_{t-1} + W^2 x_t)$

- $h_t = h_{t-1} + u_t$

- then, $\left\| \frac{\partial h_t}{\partial h_{t-1}} \right\| = 1 + \left\| \frac{\partial u_t}{\partial h_{t-1}} \right\| \geq 1$

- On the other hand

- $h_t = h_{t-1} + u_t = h_{t-2} + u_{t-1} + u_t = \dots$



A better “solution” (cont.)

- Interpolate between old state and new state (“choosing to **forget**”)
 - $f_t = \sigma(W^f x_t + U^f h_{t-1})$
 - $h_t = f_t \odot h_{t-1} + (1 - f_t) \odot u_t$
- Introduce a separate **input gate** i_t
 - $i_t = \sigma(W^i x_t + U^i h_{t-1})$
 - $h_t = f_t \odot h_{t-1} + i_t \odot u_t$
- Selectively expose memory cell c_t with an **output gate** o_t
 - $o_t = \sigma(W^o x_t + U^o h_{t-1})$
 - $c_t = f_t \odot c_{t-1} + i_t \odot u_t$
 - $h_t = o_t \odot \tanh(c_t)$

Long Short-Term Memory (LSTM)

$$u_t = \tanh(W h_{t-1} + V x_t)$$

$$f_t = \text{sigmoid}(W_f h_{t-1} + V_f x_t)$$

$$i_t = \text{sigmoid}(W_i h_{t-1} + V_i x_t)$$

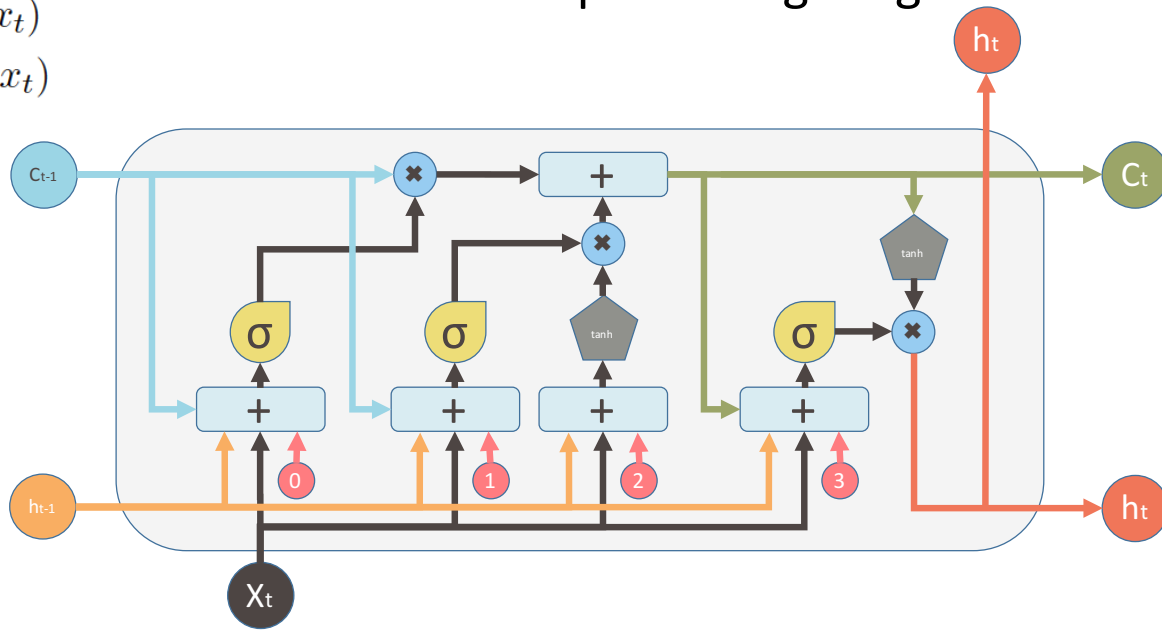
$$o_t = \text{sigmoid}(W_o h_{t-1} + V_o x_t)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot u_t$$

$$h_t = o_t \odot \tanh(c_t)$$

$$y_t = U h_t$$

- Hochreiter & Schmidhuber, 1997
- LSTM = additive updates + gating



Gated Recurrent Unites, GRU (Cho et al. 2014)

- Main ideas
 - Keep around memories to capture long distance dependencies
 - Allow error messages to flow at different strengths depending on the inputs
- Update gate
 - Based on current input and hidden state
 - $z_t = \sigma(W^z x_t + U^z h_{t-1})$
- Reset gate
 - Similarly but with different weights
 - $r_t = \sigma(W^r x_t + U^r h_{t-1})$

GRU

- New memory content

- $\tilde{h}_t = \tanh(Wx_t + r_t \odot Uh_{t-1})$

- Final memory at time step combines current and previous time steps

- $h_t = z_t \odot h_{t-1} + (1 - z_t) \odot \tilde{h}$


- Update gate z controls how much of past state should matter now

- If z closed to 1, then we can copy information in that unit through many time steps \rightarrow less vanishing gradient!

- If reset gate r unit is close to 0, then this ignores previous memory and only stores the new input information \rightarrow allows model to drop information that is irrelevant in the future

- Units with long term dependencies have active update gates z

- Units with short-term dependencies often have reset gates r very active




LSTM vs. GRU


- No clear winner!
- Tuning hyperparameters like layer size is probably more important than picking the ideal architecture
- GRUs have fewer parameters and thus may train a bit faster or need less data to generalize
- If you have enough data, the greater expressive power of LSTMs may lead to better results.

More RNNs

- Bidirectional RNN




- Deep Bidirectional RNN




Tree-LSTMs

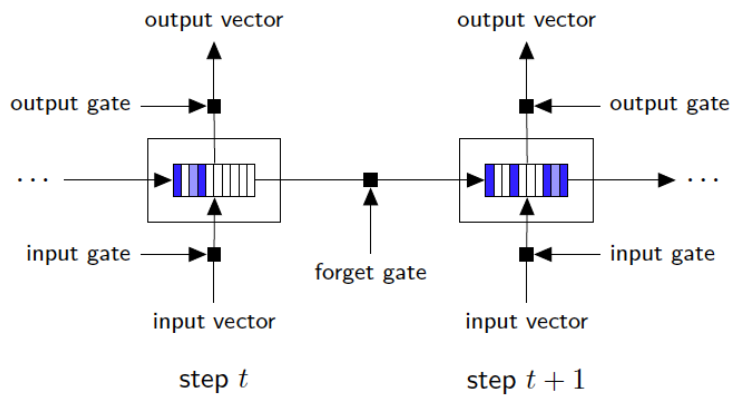
- Traditional Sequential Composition




- Tree-Structured Composition



Tree-Structured LSTMs



- Standard LSTMs
 - Each node has one child



- Tree-LSTMs
 - Accept multiple children


Conditioned Language Models

- Not just generate text, generate text according to some specification

<u>Input X</u>	<u>Output Y (Text)</u>	<u>Task</u>
Structured Data	NL Description	NL Generation
English	Japanese	Translation
Document	Short Description	Summarization
Utterance	Response	Response Generation
Image	Text	Image Captioning
Speech	Transcript	Speech Recognition

Neural Machine Translation

- The Vauquois Diagram (1968)




Neural Machine Translation


- RNN trained end-to-end: encoder-decoder

“You can’t cram the meaning of a whole %&!\$ing sentence into a single \$&!*ing vector!” — Ray Mooney

Encoder:

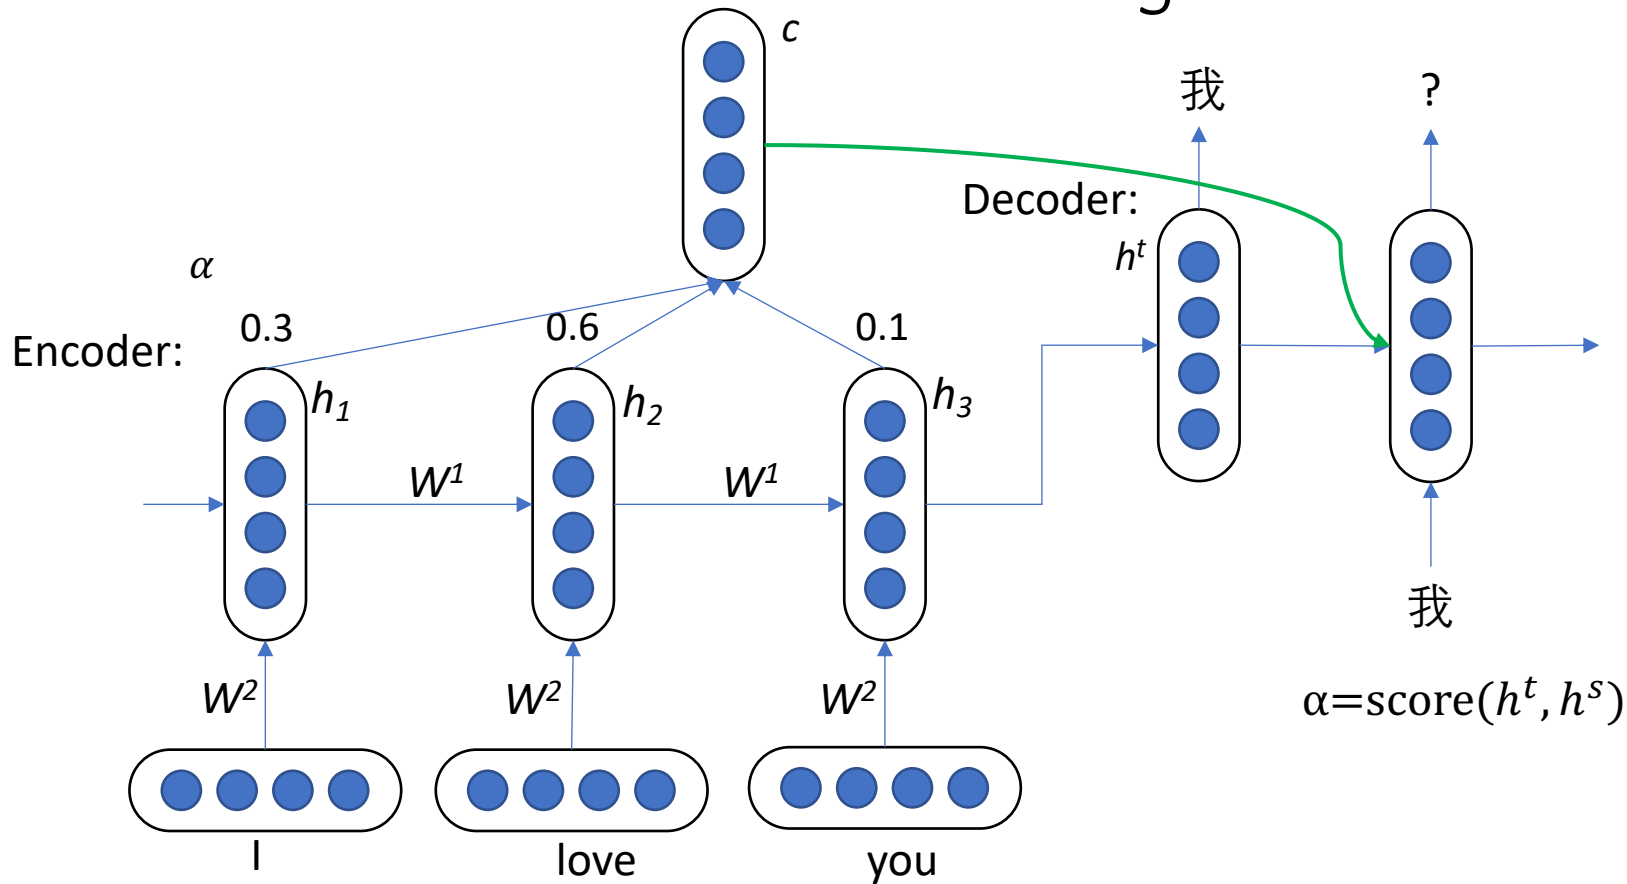


Decoder:



This needs to capture the entire sentence!

Attention Mechanism – *Scoring*



Attention Score Functions

- \mathbf{q} is the query and \mathbf{k} is the key
- **Multi-layer Perceptron** (Bahdanau et al. 2015)
 - $a(\mathbf{q}, \mathbf{k}) = \mathbf{w}_2^T \tanh(W_1[\mathbf{q}; \mathbf{k}])$
 - Flexible, often very good with large data
- **Bilinear** (Luong et al. 2015)
 - $a(\mathbf{q}, \mathbf{k}) = \mathbf{q}^T W \mathbf{k}$
- **Dot Product** (Luong et al. 2015)
 - $a(\mathbf{q}, \mathbf{k}) = \mathbf{q}^T \mathbf{k}$
 - No parameters! But requires sizes to be the same.
- **Scaled Dot Product** (Vaswani et al. 2017)
 - Problem: scale of dot product increases as dimensions get larger
 - Fix: scale by size of the vector
 - $a(\mathbf{q}, \mathbf{k}) = \frac{\mathbf{q}^T \mathbf{k}}{\sqrt{|\mathbf{k}|}}$

Results

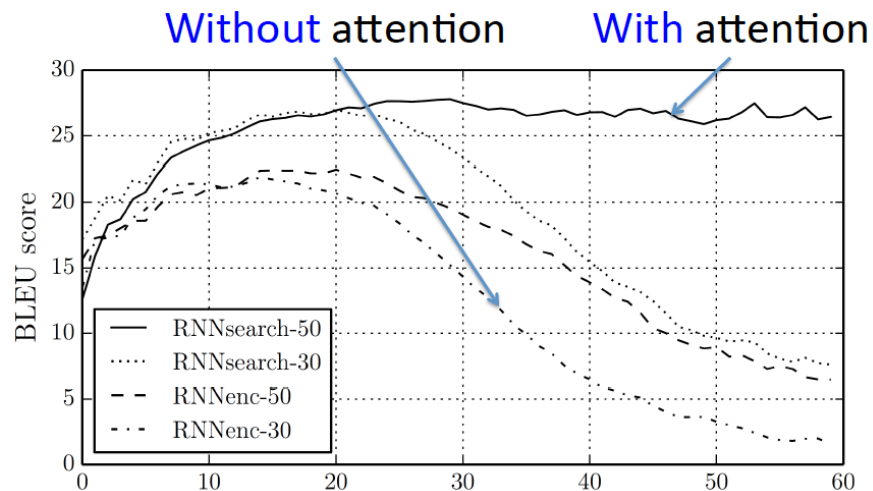



Image Caption Generation (Xu et al., 2015)

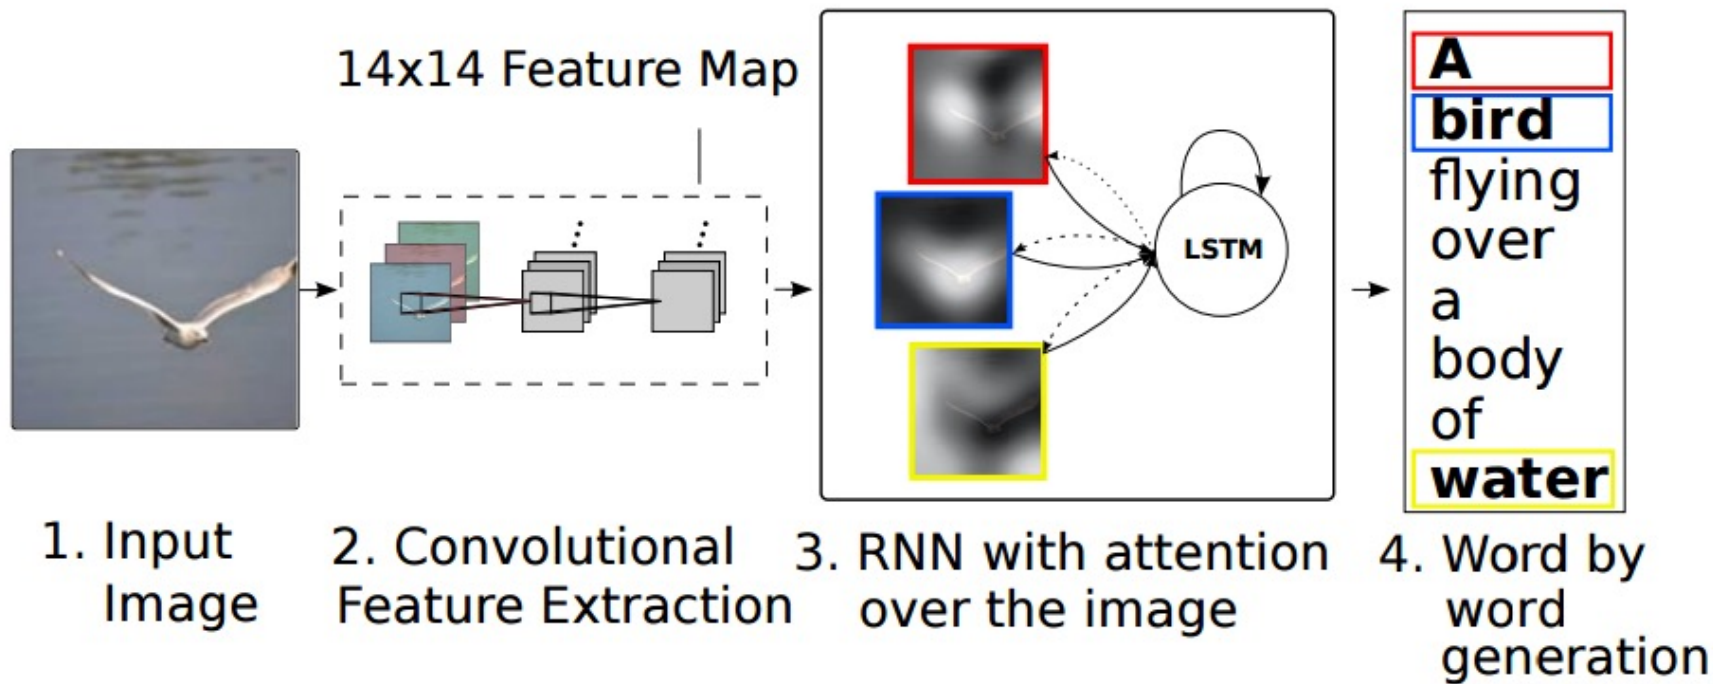


Image Caption Generation (Xu et al., 2015)






Image Caption Generation (Xu et al., 2015)




A woman is throwing a frisbee in a park.



A dog is standing on a hardwood floor.



A little girl sitting on a bed with a teddy bear.




A group of people sitting on a boat in the water.



Image Caption Generation (Xu et al., 2015)



A large white bird standing in a forest.



A woman holding a clock in her hand.



A man wearing a hat and
a hat on a skateboard.



A person is standing on a beach
with a surfboard.

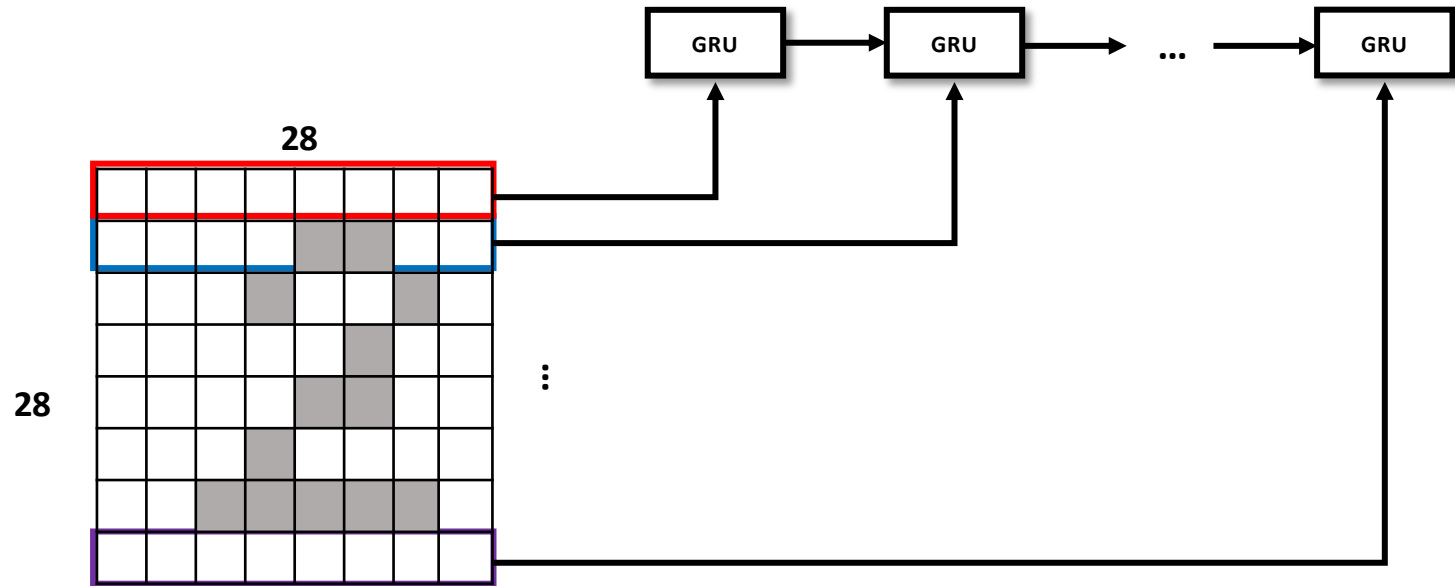


A woman is sitting at a table
with a large pizza.

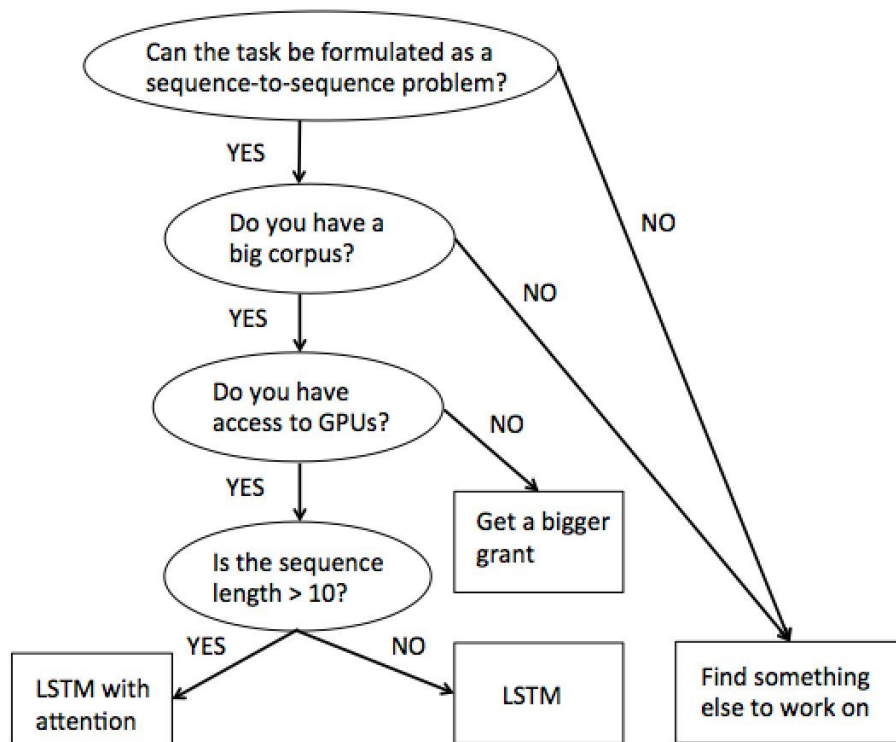


A man is talking on his cell phone
while another man watches.

RNN for MNIST



How to do NLP research now?



Transformers

What's Transformers

- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, Illia Polosukhin. **Attention Is All You Need.** NIPS 2017.
- Definition
 - Any architecture designed to process a connected set of units—such as the tokens in a sequence or the pixels in an image—where the only interaction between units is through **self-attention**.
- Key component
 - Self-attention




Self-attention

- A sequence-to-sequence operation
 - Input vectors: $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_t$
 - Output vectors: $\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_t$
 - The vectors all have dimension k
- $\mathbf{y}_i = \sum_j w_{ij} \mathbf{x}_j$, where j indexes over the whole sequence
- w_{ij} is derived from a function over \mathbf{x}_i and \mathbf{x}_j
- The simplest option for this function is the dot product:
 - $w'_{ij} = \mathbf{x}_i^T \mathbf{x}_j$
 - Then apply a softmax to map the values to $[0, 1]$: $w_{ij} = \frac{\exp w'_{ij}}{\sum_j \exp w'_{ij}}$




A visual illustration of basic self-attention



Self-attention is the only operation in Transformers that propagates information between vectors

Understanding why self-attention works


- For example, movie recommendation
- Using features to represent movies and users
 - Movies: how much romance there is in the movie, and how much action
 - Users: how much they enjoy romantic movies and how much they enjoy action-based movies
- The dot product between the two feature vectors would give you a score for how well a movie match what a user enjoys




The signs of a feature are meaningful.

How to assign features to movies and users?

- Ask users for a small number of movies that they like (training data)
- Optimize the user and movie features so that their dot product matches the data



Self-attention for a sequence of words



Self-attention in action


```
1  import torch
2  import torch.nn.functional as F
3
4  x = torch.rand(3, 4, 5)
5
6  # torch.bmm is a batched matrix multiplication. It
7  # applies matrix multiplication over batches of
8  # matrices.
9
10 raw_weights = torch.bmm(x, x.transpose(1, 2))
11
12 weights = F.softmax(raw_weights, dim=2)
13
14 y = torch.bmm(weights, x)
15
16 print(y)
```


Shortcomings of basic self-attention

- A vector plays three roles
 - Vectors \mathbf{x}_i and \mathbf{x}_j to establish the weights
 - Vectors \mathbf{x}_j to establish the score
- Dot product grows with the embedding dimension k
 - Softmax function can be sensitive to very large input values, which kill the gradient and slow down learning
- A word can mean different things to different neighbors
- The results are position-independent

Queries, keys and values

- A vector plays three roles
 - Using three $k \times k$ matrices to map \mathbf{x}_i into three vectors to play different roles
- These roles are often called the **query**, the **key** and the **value**
 - $\mathbf{q}_i = \mathbf{W}_q \mathbf{x}_i$, $\mathbf{k}_i = \mathbf{W}_k \mathbf{x}_i$, $\mathbf{v}_i = \mathbf{W}_v \mathbf{x}_i$
 - $w_{ij}' = \mathbf{q}_i^T \mathbf{k}_j$
 - $w_{ij} = \text{softmax}(w_{ij}')$
 - $\mathbf{y}_i = \sum_j w_{ij} \mathbf{v}_j$




Scaling the dot product

- $w'_{ij} = \frac{q_i^T k_j}{\sqrt{k}}$, where k is the dimension of vector x_i
- Why \sqrt{k} ?
 - Imagine a vector in \mathbb{R}^k with values all c . Its Euclidean length is $\sqrt{k}c$
 - Dividing out the amount by which the increase in dimension increases the length of the average vectors.


Multi-head attention

- A word can mean different things to different neighbors
 - For example, “the cat **walks** on the street”
- Combining several self-attention mechanisms, each with different matrices \mathbf{W}_q^r , \mathbf{W}_k^r , \mathbf{W}_v^r (r is the index of a head)
- For input \mathbf{x}_i each attention head produces a different output vector \mathbf{y}_i^r . We concatenate these, and pass them through a linear transformation to reduce the dimension back to k



Multi-head attention seems ensemble of attention

Using the positions



- Position embeddings

- Embedding the positions like we did the words
- For example
 - $\mathbf{v}_{12} = [0.4259, 0.1208, 0.5587]$, $\mathbf{v}_{25} = [0.7838, 0.4896, 0.5110]$

- Position encodings

- Using a function $f: \mathbb{N} \rightarrow \mathbb{R}^k$ to map a positions to a real valued vector

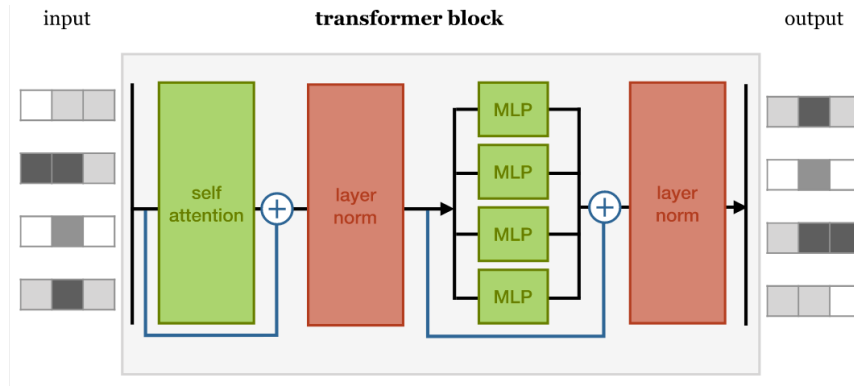
$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{model}})$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{model}})$$

	<	—	emb_dim	—	>
The	$\sin\left(\frac{0}{10000^{\frac{0}{emb_dim}}}\right)$	$\cos\left(\frac{0}{10000^{\frac{0}{emb_dim}}}\right)$	$\sin\left(\frac{0}{10000^{\frac{2}{emb_dim}}}\right)$	$\cos\left(\frac{0}{10000^{\frac{2}{emb_dim}}}\right)$	\dots
daughter	$\sin\left(\frac{1}{10000^{\frac{0}{emb_dim}}}\right)$	$\cos\left(\frac{1}{10000^{\frac{0}{emb_dim}}}\right)$	$\sin\left(\frac{1}{10000^{\frac{2}{emb_dim}}}\right)$	$\cos\left(\frac{1}{10000^{\frac{2}{emb_dim}}}\right)$	\dots
had	$\sin\left(\frac{2}{10000^{\frac{0}{emb_dim}}}\right)$	$\cos\left(\frac{2}{10000^{\frac{0}{emb_dim}}}\right)$	$\sin\left(\frac{2}{10000^{\frac{2}{emb_dim}}}\right)$	$\cos\left(\frac{2}{10000^{\frac{2}{emb_dim}}}\right)$	\dots
a	$\sin\left(\frac{3}{10000^{\frac{0}{emb_dim}}}\right)$	$\cos\left(\frac{3}{10000^{\frac{0}{emb_dim}}}\right)$	$\sin\left(\frac{3}{10000^{\frac{2}{emb_dim}}}\right)$	$\cos\left(\frac{3}{10000^{\frac{2}{emb_dim}}}\right)$	\dots
nice	$\sin\left(\frac{4}{10000^{\frac{0}{emb_dim}}}\right)$	$\cos\left(\frac{4}{10000^{\frac{0}{emb_dim}}}\right)$	$\sin\left(\frac{4}{10000^{\frac{2}{emb_dim}}}\right)$	$\cos\left(\frac{4}{10000^{\frac{2}{emb_dim}}}\right)$	\dots
umbrella	$\sin\left(\frac{5}{10000^{\frac{0}{emb_dim}}}\right)$	$\cos\left(\frac{5}{10000^{\frac{0}{emb_dim}}}\right)$	$\sin\left(\frac{5}{10000^{\frac{2}{emb_dim}}}\right)$	$\cos\left(\frac{5}{10000^{\frac{2}{emb_dim}}}\right)$	\dots


Transformer block

- Self-attention + non-linear
 - Combine **self-attention** with a **local feedforward**




- **Normalization** and **residual connections** are standard tricks used to help DNNs train faster and more accurately

Batch Normalization

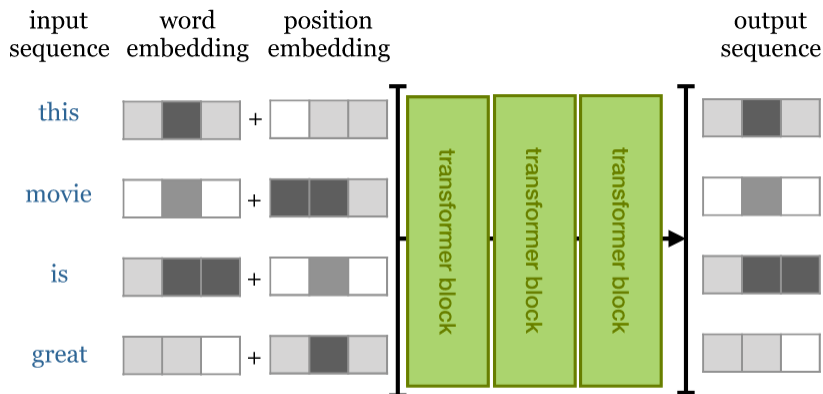


Layer Normalization



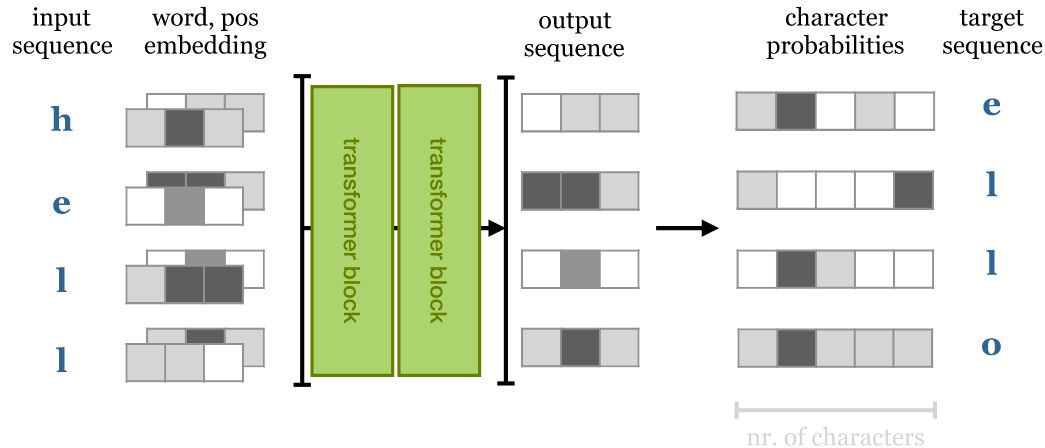
Transformers for text classification

- Transformer
 - A chain of transformer blocks
- Text classification
 - Apply global average pooling to the final output sequence




Transformers for text generation

- Autoregressive model
- For example
 - Predict the next character in a sequence



Transformers for text generation

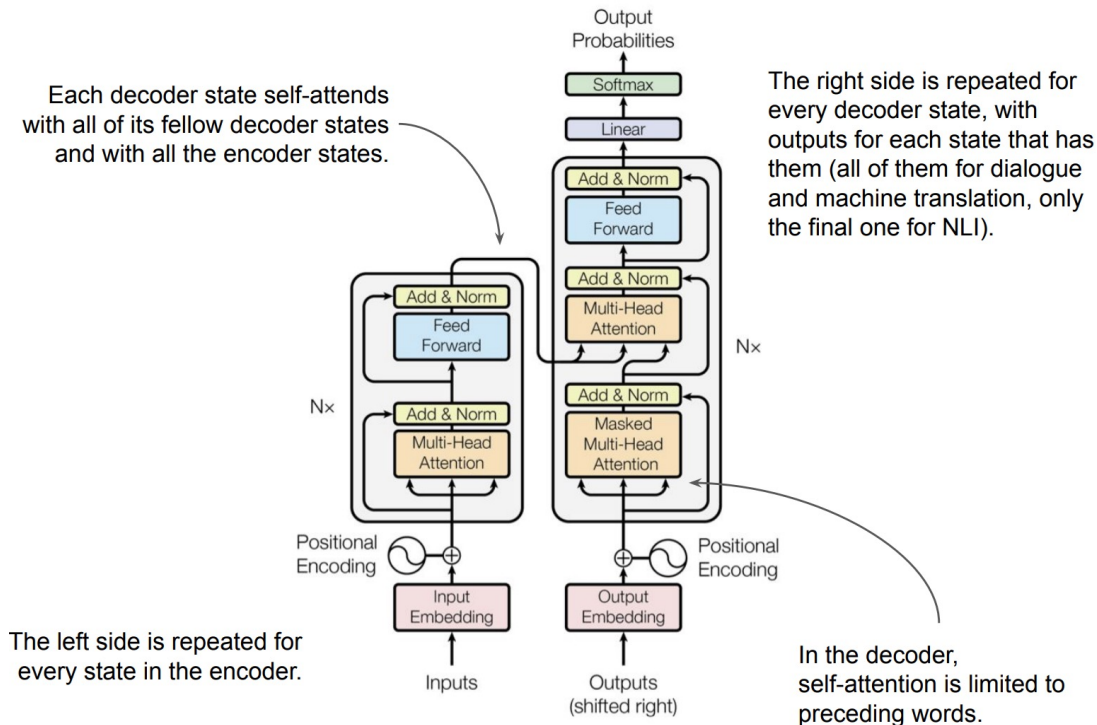
- Output token should only depend only on previous inputs
- With a transformer, the output depends on the entire input sequence
- Applying a mask to the matrix of dot products, before the softmax is applied
 - Set the masked out elements to -inf





Transformers for encoder-decoder

- The decoder generates the output sentence word for word based both on the input words and the words it has already generated

If you can use seq2seq, you can use transformer.



Attention Visualization





The encoder self-attention distribution for the word "it" from the 5th to the 6th layer of a Transformer trained on English to French translation (one of eight attention heads).

<https://ai.googleblog.com/2017/08/transformer-novel-neural-network.html>

Advantages of Transformers

- Compared to RNN, Transformers have no recurrent connections
- So the whole model can be computed in a very efficient
- Compared to CNN, Transformers can model dependencies over the whole range of the input sequence



Disadvantages of Transformers

- The model is too big to be fitted in a GPU
- Cannot handle too long sequences
- How to solve the problem?
 - Half precision
 - Gradient accumulation
 - Gradient checkpointing
 - <https://medium.com/huggingface/training-larger-batches-practical-tips-on-1-gpu-multi-gpu-distributed-setups-ec88c3e51255>

Summary: RNNs and Transformers

RNNs

- Recurrent Neural Networks
- The vanishing gradient problem
- Long Short-Term Memory (LSTM)
- Gated Recurrent Unites (GRU)
- Attention mechanism
- Can be used in almost all applications

Transformers

- Basic self-attention
- Improvements
 - Query, key and value transformation
 - Scaling the dot product
 - Multi-head attention
 - Positional embeddings and encodings
- Transformers are extremely generic
- The current performance limit is purely in the hardware

参考

- 《自然语言处理：基于预训练模型的方法》
 - 出版社：电子工业出版社
 - 作者：车万翔，郭江，崔一鸣 著；刘挺 主审
 - 书号：ISBN 978-7-121-41512-8
 - 出版时间：2021.7
- 网购链接
 - <https://item.jd.com/13344628.html>
- 书中代码
 - <https://github.com/HIT-SCIR/plm-nlp-code>

